


Shared Memory Programming with Pthreads & OpenMP

Dilum Bandara
Dilum.Bandara@uom.lk

Slides extended from
An Introduction to Parallel Programming by
Peter Pacheco

Shared Memory System



POSIX[®] Threads

- Also known as Pthreads
- Standard for Unix-like operating systems
- Library that can be linked with C programs
- Specifies an API for multi-threaded programming

Hello World!

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
```

Declares various Pthreads functions, constants, types, etc.



```
/* Global variable: accessible to all threads */
int thread_count;
```

```
void *Hello(void* rank); /* Thread function */
```

```
int main(int argc, char* argv[]) {
    long          thread; /* Use long in case of a 64-bit system */
    pthread_t* thread_handles;
```

```
/* Get number of threads from command line */
thread_count = strtol(argv[1], NULL, 10);
```

```
thread_handles = malloc (thread_count*sizeof(pthread_t));
```

Hello World! (Cont.)

```
for (thread = 0; thread < thread_count; thread++)  
    pthread_create(&thread_handles[thread], NULL,  
        Hello, (void*) thread);  
  
printf("Hello from the main thread\n");  
  
for (thread = 0; thread < thread_count; thread++)  
    pthread_join(thread_handles[thread], NULL);  
  
free(thread_handles);  
return 0;  
} /* main */
```


Hello World! (Cont.)

```
void *Hello(void* rank) {  
    long my_rank = (long) rank;  /* Use long in case of 64-bit system */  
  
    printf("Hello from thread %ld of %d\n", my_rank, thread_count);  
  
    return NULL;  
} /* Hello */
```

Compiling a Pthread program

```
gcc -g -Wall -o pth_hello pth_hello.c -lpthread
```

Link Pthreads library



Running a Pthreads program

. /pth_hello <number of threads>

. /pth_hello 1

Hello from the main thread

Hello from thread 0 of 1

. /pth_hello 4

Hello from the main thread


Hello from thread 0 of 4

Hello from thread 3 of 4

Hello from thread 2 of 4

Hello from thread 1 of 4

Running the Threads



Main thread forks & joins 2 threads

Global Variables

- Can introduce subtle & confusing bugs!
- Use them only when they are essential
 - Shared variables



Starting Threads

pthread.h

One object for
each thread

pthread_t

We ignore return value
from pthread_create

```
int pthread_create (  
    pthread_t* thread_p,           /* out */  
    const pthread_attr_t* attr_p,  /* in */  
    void* (*start_routine) (void), /* in */  
    void* arg_p);                  /* in */
```

Function Started by pthread_create

- Function start by pthread_create should have following prototype
`void* thread_function (void* args_p) ;`
- Void* can be cast to any pointer type in C
 - So args_p can point to a list containing one or more values needed by thread_function
- Similarly, return value of thread_function can point to a list of one or more values

Stopping Threads

- Single call to `pthread_join` will wait for thread associated with `pthread_t` object to complete
 - Suspend execution of calling thread until target thread terminates, unless it has already terminated
 - Call `pthread_join` once for each thread

```
int pthread_join(  
    pthread_t* thread          /* in */,  
    void** ret_val_p          /* out */ );
```

a_{00}	a_{01}	\cdots	$a_{0,n-1}$
a_{10}	a_{11}	\cdots	$a_{1,n-1}$
\vdots	\vdots		\vdots
a_{i0}	a_{i1}	\cdots	$a_{i,n-1}$
\vdots	\vdots		\vdots
$a_{m-1,0}$	$a_{m-1,1}$	\cdots	$a_{m-1,n-1}$

x_0
 x_1
 \vdots
 x_{n-1}

=

y_0
y_1
\vdots
$y_i = a_{i0}x_0 + a_{i1}x_1 + \cdots a_{i,n-1}x_{n-1}$
\vdots
y_{m-1}

Matrix-Vector Multiplication in Pthreads

Serial Pseudo-code

$$y_i = \sum_{j=0}^{n-1} a_{ij}x_j$$

```
/* For each row of A */
for (i = 0; i < m; i++) {
    y[i] = 0.0;
    /* For each element of the row and each element of x */
    for (j = 0; j < n; j++)
        y[i] += A[i][j]* x[j];
}
```

Using 3 Pthreads

- Assign each row to a separate thread
- Suppose 6x6 matrix & 3 threads

Thread	Components of y
0	y[0], y[1]
1	y[2], y[3]
2	y[4], y[5]

Thread 0

```
y[0] = 0.0;  
for (j = 0; j < n; j++)  
    y[0] += A[0][j]* x[j];
```

General case

```
y[i] = 0.0;  
for (j = 0; j < n; j++)  
    y[i] += A[i][j]*x[j];
```


Pthreads Matrix-Vector Multiplication

```
void *Pth_mat_vect(void* rank) {  
    long my_rank = (long) rank;  
    int i, j;  
    int local_m = m/thread_count;  
    int my_first_row = my_rank*local_m;  
    int my_last_row = (my_rank+1)*local_m - 1;  
  
    for (i = my_first_row; i <= my_last_row; i++) {  
        y[i] = 0.0;  
        for (j = 0; j < n; j++)  
            y[i] += A[i][j]*x[j];  
    }  
  
    return NULL;  
} /* Pth_mat_vect */
```

Estimating π

$$\pi = 4 \left(1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \cdots + (-1)^n \frac{1}{2n+1} + \cdots \right)$$

```
double factor = 1.0;
double sum = 0.0;
for (i = 0; i < n; i++, factor = -factor) {
    sum += factor/(2*i+1);
}
pi = 4.0*sum;
```

Thread Function for Computing π

```
void* Thread_sum(void* rank) {
    long my_rank = (long) rank;
    double factor;
    long long i;
    long long my_n = n/thread_count;
    long long my_first_i = my_n*my_rank;
    long long my_last_i = my_first_i + my_n;

    if (my_first_i % 2 == 0)  /* my_first_i is even */
        factor = 1.0;
    else /* my_first_i is odd */
        factor = -1.0;

    for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
        sum += factor/(2*i+1);
    }

    return NULL;
} /* Thread_sum */
```

Using a dual core processor

	n			
	10^5	10^6	10^7	10^8
π	3.14159	3.141593	3.1415927	3.14159265
1 Thread	3.14158	3.141592	3.1415926	3.14159264
2 Threads	3.14158	3.141480	3.1413692	3.14164686

As we increase n , estimate with 1 thread gets better & better

2 thread case produce different answers in different runs

Why?

Pthreads Global Sum with Busy-Waiting

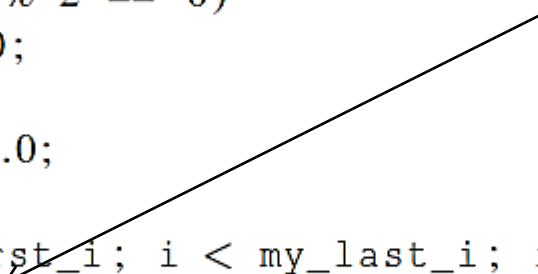
```
void* Thread_sum(void* rank) {
    long my_rank = (long) rank;
    double factor;
    long long i;
    long long my_n = n/thread_count;
    long long my_first_i = my_n*my_rank;
    long long my_last_i = my_first_i + my_n;

    if (my_first_i % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;

    for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
        while (flag != my_rank);
        sum += factor/(2*i+1);
        flag = (flag+1) % thread_count;
    }

    return NULL;
} /* Thread_sum */
```

Shared variable



Mutexes

- Make sure only 1 thread in critical region
- Pthreads standard includes a special type for mutexes: `pthread_mutex_t`

```
int pthread_mutex_init(  
    pthread_mutex_t*      mutex_p      /* out */  
    const pthread_mutexattr_t* attr_p    /* in  */);
```

Mutexes

- Lock

- To gain access to a critical section

```
int pthread_mutex_lock(pthread_mutex_t* mutex_p /* in/out */);
```

- Unlock

- When a thread is finished executing code in a critical section

```
int pthread_mutex_unlock(pthread_mutex_t* mutex_p /* in/out */);
```

- Termination

- When a program finishes using a mutex

```
int pthread_mutex_destroy(pthread_mutex_t* mutex_p /* in/out */);
```

Global Sum Function Using a Mutex

```
void* Thread_sum(void* rank) {  
    long my_rank = (long) rank;  
    double factor;  
    long long i;  
    long long my_n = n/thread_count;  
    long long my_first_i = my_n*my_rank;  
    long long my_last_i = my_first_i + my_n;  
    double my_sum = 0.0;  
  
    if (my_first_i % 2 == 0)  
        factor = 1.0;  
    else  
        factor = -1.0;
```


Global Sum Function Using a Mutex (Cont.)

```
for (i = my_first_i; i < my_last_i; i++, factor = -factor) {  
    my_sum += factor/(2*i+1);  
}  
pthread_mutex_lock(&mutex);  
sum += my_sum;  
pthread_mutex_unlock(&mutex);  
  
return NULL;  
} /* Thread_sum */
```

Busy-Waiting vs. Mutex

Threads	Busy-Wait	Mutex
1	2.90	2.90
2	1.45	1.45
4	0.73	0.73
8	0.38	0.38
16	0.50	0.38
32	0.80	0.40
64	3.56	0.38

$$\frac{T_{\text{serial}}}{T_{\text{parallel}}} \approx \text{thread_count}$$

Run-times (in seconds) of π programs using $n = 108$ terms on a system with 2x4-core processors

Semaphores

```
#include <semaphore.h>
```

← Semaphores are not part of Pthreads;
you need to add this


```
int sem_init(  
    sem_t*      semaphore_p    /* out */,  
    int         shared         /* in */,  
    unsigned    initial_val    /* in */);
```

```
int sem_destroy(sem_t*      semaphore_p /* in/out */);  
int sem_post(sem_t*        semaphore_p /* in/out */);  
int sem_wait(sem_t*        semaphore_p /* in/out */);
```

Read-Write Locks

- While controlling access to a large, shared data structure
- Example
 - Suppose shared data structure is a sorted linked list of ints, & operations of interest are Member, Insert, & Delete

Linked Lists



```
struct list_node_s {  
    int data;  
    struct list_node_s* next;  
}
```


Linked List Membership

```
int  Member(int value, struct list_node_s* head_p) {
    struct list_node_s* curr_p = head_p;

    while (curr_p != NULL && curr_p->data < value)
        curr_p = curr_p->next;

    if (curr_p == NULL || curr_p->data > value) {
        return 0;
    } else {
        return 1;
    }
} /* Member */
```

Inserting New Node Into a List



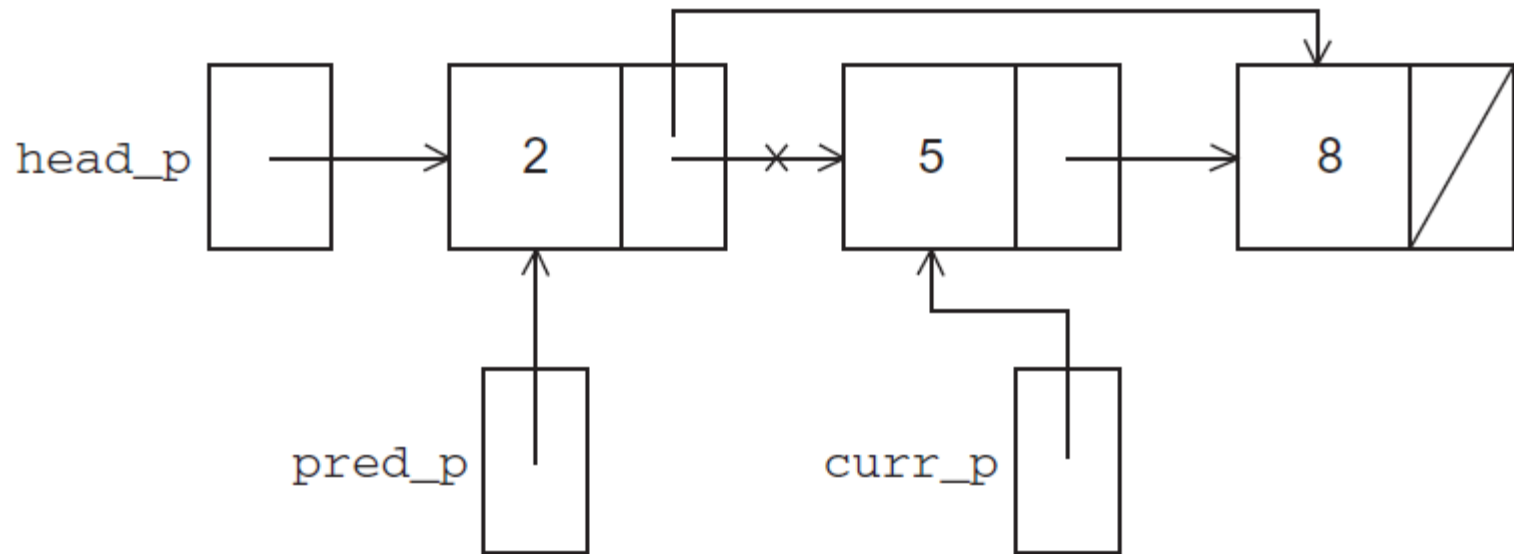
Inserting New Node Into a List (Cont.)

```
int Insert(int value, struct list_node_s** head_pp) {
    struct list_node_s* curr_p = *head_pp;
    struct list_node_s* pred_p = NULL;
    struct list_node_s* temp_p;

    while (curr_p != NULL && curr_p->data < value) {
        pred_p = curr_p;
        curr_p = curr_p->next;
    }

    if (curr_p == NULL || curr_p->data > value) {
        temp_p = malloc(sizeof(struct list_node_s));
        temp_p->data = value;
        temp_p->next = curr_p;
        if (pred_p == NULL) /* New first node */
            *head_pp = temp_p;
        else
            pred_p->next = temp_p;
        return 1;
    } else { /* Value already in list */
        return 0;
    }
} /* Insert */
```


Deleting a Node From a Linked List



Deleting a Node From a Linked List (Cont.)

```
int Delete(int value, struct list_node_s** head_pp) {
    struct list_node_s* curr_p = *head_pp;
    struct list_node_s* pred_p = NULL;

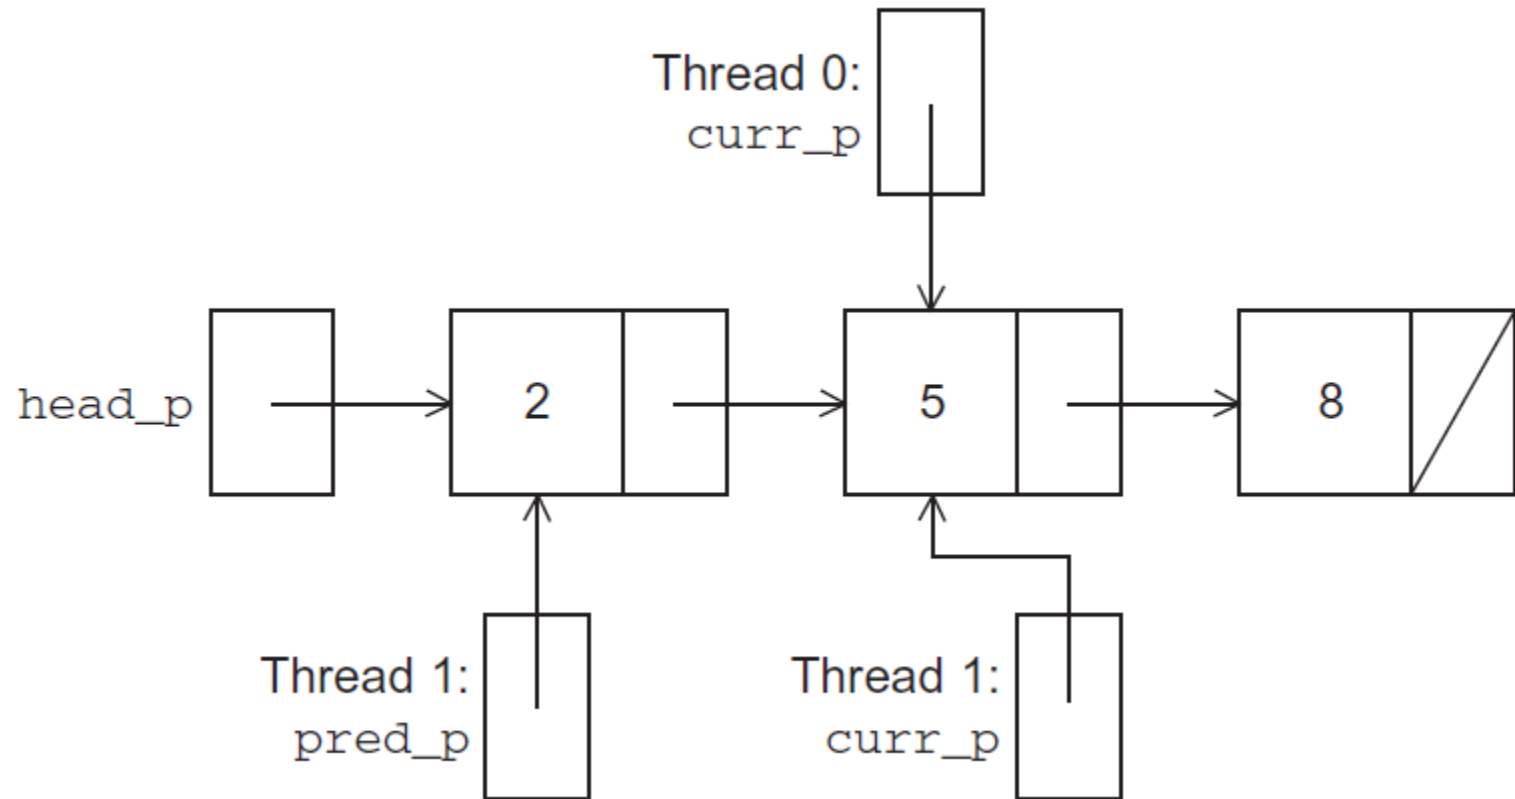
    while (curr_p != NULL && curr_p->data < value) {
        pred_p = curr_p;
        curr_p = curr_p->next;
    }

    if (curr_p != NULL && curr_p->data == value) {
        if (pred_p == NULL) { /* Deleting first node in list */
            *head_pp = curr_p->next;
            free(curr_p);
        } else {
            pred_p->next = curr_p->next;
            free(curr_p);
        }
        return 1;
    } else { /* Value isn't in list */
        return 0;
    }
} /* Delete */
```

Multi-Threaded Linked List

- To share access to the list, we can define `head_p` to be a global variable
 - This will simplify function headers for `Member`, `Insert`, & `Delete`
 - Because we won't need to pass in either `head_p` or a pointer to `head_p`: we'll only need to pass in the value of interest

Simultaneous Access by 2 Threads



Solution #1

- Simply lock the list any time that a thread attempts to access it
- Call to each of the 3 functions can be protected by a mutex

```
Pthread_mutex_lock(&list_mutex);  
Member(value);  
Pthread_mutex_unlock(&list_mutex);
```

In place of calling Member(value).

Issues

- Serializing access to the list
- If vast majority of our operations are calls to **Member**
 - We fail to exploit opportunity for parallelism
- If most of our operations are calls to **Insert & Delete**
 - This may be the best solution

Solution #2

- Instead of locking entire list, we could try to lock individual nodes
- A “finer-grained” approach

```
struct list_node_s {  
    int data;  
    struct list_node_s* next;  
    pthread_mutex_t mutex;  
}
```

Issues

- Much more complex than original Member function
- Much slower
 - Because each time a node is accessed, a mutex must be locked & unlocked
 - Addition of a mutex field to each node substantially increase memory needed for the list

Pthreads Read-Write Locks

- Neither multi-threaded linked lists exploits potential for simultaneous access to any node by threads that are executing Member
 - 1st solution only allows 1 thread to access the entire list at any instant
 - 2nd only allows 1 thread to access any given node at any instant
- Read-write lock is somewhat like a mutex except that it provides 2 lock functions
- 1st locks the read-write lock for reading
- 2nd locks it for writing

Pthreads Read-Write Locks (Cont.)

- Multiple threads can simultaneously obtain lock by calling read-lock function
- While only 1 thread can obtain lock by calling write-lock function
- Thus
 - If any thread owns lock for reading, any thread that wants to obtain a lock for writing will be blocked
 - If any thread owns lock for writing, any threads that want to obtain lock for reading or writing will be blocked

Protecting Our Linked List Functions

```
pthread_rwlock_rdlock(&rwlock);
Member(value);
pthread_rwlock_unlock(&rwlock);
. . .
pthread_rwlock_wrlock(&rwlock);
Insert(value);
pthread_rwlock_unlock(&rwlock);
. . .
pthread_rwlock_wrlock(&rwlock);
Delete(value);
pthread_rwlock_unlock(&rwlock);
```

Linked List Performance

100,000 ops/thread
99.9% Member
0.05% Insert
0.05% Delete

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	0.213	0.123	0.098	0.115
One Mutex for Entire List	0.211	0.450	0.385	0.457
One Mutex per Node	1.680	5.700	3.450	2.700

100,000 ops/thread
80% Member
10% Insert
10% Delete

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	2.48	4.97	4.69	4.71
One Mutex for Entire List	2.50	5.13	5.04	5.11
One Mutex per Node	12.00	29.60	17.00	12.00

OpenMP

OpenMP

- High-level API for shared-memory parallel programming
 - MP = multiprocessing
- Use Pragmas
 - Special preprocessor instructions
 - #pragma
 - Typically added to support behaviors that aren't part of the basic C specification
 - Compilers that don't support pragmas ignore them

```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);

} /* Hello */

```

Compiling & Running

```
gcc -g -Wall -fopenmp -o omp_hello omp_hello . c
```

```
./ omp_hello 4
```

running with 4 threads

compiling

Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4


possible
outcomes

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4

Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4

OpenMp pragmas

- **# pragma omp parallel**
 - Most basic parallel directive
 - Original thread is called **master**
 - Additional threads are called **slaves**
 - Original thread & new threads called a **team**



Clause

- Text that modifies a directive
- *num_threads* clause can be added to a parallel directive
- Allows programmer to specify no of threads that should execute following block

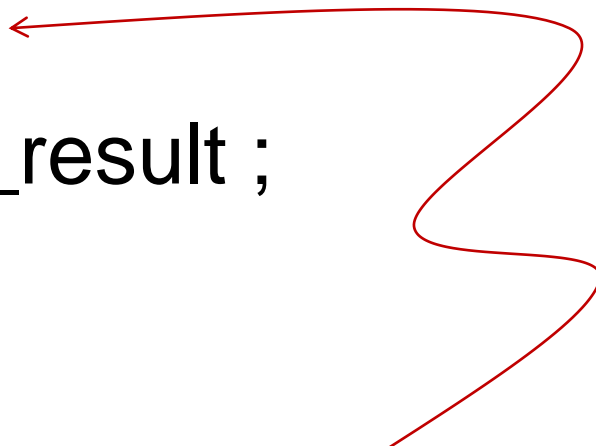
```
# pragma omp parallel num_threads ( thread_count )
```

Be Aware...

- There may be system-defined limitations on number of threads that a program can start
- OpenMP standard doesn't guarantee that this will actually start *thread_count* threads
- Most current systems can start hundreds or even 1,000s of threads
- Unless we're trying to start a lot of threads, we will almost always get desired no of threads



Mutual Exclusion

```
# pragma omp critical  
    global_result += my_result ;
```



only 1 thread can execute following
structured block at a time


Trapezoidal Rule



Serial algorithm

```
/* Input:  a, b, n */  
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++) {  
    x_i = a + i*h;  
    approx += f(x_i);  
}  
approx = h*approx;
```

Assignment of Trapezoids to Threads



```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Trap(double a, double b, int n, double* global_result_p);

int main(int argc, char* argv[]) {
    double    global_result = 0.0;  /* Store result in global_result */
    double    a, b;                 /* Left and right endpoints      */
    int       n;                    /* Total number of trapezoids    */
    int       thread_count;

    thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);
    # pragma omp parallel num_threads(thread_count)
    Trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n",
        a, b, global_result);
    return 0;
} /* main */

```

```

void Trap(double a, double b, int n, double* global_result_p) {
    double  h, x, my_result;
    double  local_a, local_b;
    int  i, local_n;
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b-a)/n;
    local_n = n/thread_count;
    local_a = a + my_rank*local_n*h;
    local_b = local_a + local_n*h;
    my_result = (f(local_a) + f(local_b))/2.0;
    for (i = 1; i <= local_n-1; i++) {
        x = local_a + i*h;
        my_result += f(x);
    }
    my_result = my_result*h;

    # pragma omp critical
        *global_result_p += my_result;
} /* Trap */

```